Getting a New Way of Growing Adopted

Seven books on diffusion, positioning, behaviour change and pricing. One page each, read against a single problem: selling a system whose value only appears after the grower changes how they produce.

Prepared for Guillaume · August 2026

Diffusion of Innovations

Everett Rogers, 1962 · 5th edition 2003

> Adoption is not a decision about a product. It is a process of reducing uncertainty about a practice, and it runs on social proof rather than information.

Rogers spent a career on one puzzle: why demonstrably better practices spread slowly. His founding case is hybrid seed corn in Iowa, where the yield advantage was large, provable and published, and general adoption still took roughly a decade. Extension agents had the data and delivered it competently, so the data was not the bottleneck. What moved farmers was watching a neighbour plant a few acres and seeing the harvest come in. The model he builds from it has the adopter passing through knowledge, persuasion, decision, implementation and confirmation, stalling readily at any stage. Most sellers optimize for the first and assume the rest follows, which fails because the later stages turn on uncertainty about execution rather than about merit.

Relative advantage is perceived, not measured

The comparison the adopter runs is not your model against their current output. It is your model against their current output, discounted by their private confidence that they personally can execute it, minus what they lose during the transition year. A grower who fully believes a 4x gain is achievable in principle but rates their own odds at fifty percent, and expects a weaker first season while learning, is looking at a much smaller and much later number than the one on your slide. This is why a better spreadsheet almost never changes a decision: the disagreement is about their own capability, and external evidence does not speak to it.

Compatibility with existing practice and identity

An innovation that requires the adopter to become a different kind of operator meets resistance that has nothing to do with economics. Intensive production carries associations with industrial growing, and a diversified organic market gardener may reject it as a category error before evaluating a single number. Rogers found compatibility often outweighs relative advantage in explaining adoption rates, which is uncomfortable if your case rests on advantage alone. The response is framing rather than proof: present intensification as a fuller expression of what a careful grower already does, not as a move toward a system they define themselves against.

Complexity is the tax you can actually reduce

Every additional thing the adopter must learn, sequence, monitor or remember multiplies perceived risk, and perceived risk compounds rather than adds. Your approach is genuinely complex, normally a hard ceiling on adoption rate. It is also the one attribute in his list that a service can absorb on the adopter's behalf, which is what doing the daily work through technology means. Framed correctly, the largest liability becomes the offer.

Trialability decides who even considers it

Can they run it on one bay for one crop cycle without endangering the season. Innovations testable in pieces are adopted at multiple times the rate of all-or-nothing ones, because a small trial replaces secondhand belief with personal experience, the only evidence that resolves the capability question above. Trialability is a property of how you package the offer, not of the practice itself.

Observability is the engine of the second sale

Adoption spreads fastest where results are visible without needing to be explained. Yield figures in a case study are not observable in Rogers' sense. A standing crop a visiting grower can walk into and silently compare against their own is. It is why innovations with visible effects spread faster than equally valuable root-zone ones.

Early adopters and the early majority are different populations

Innovators buy on novelty and tolerate ambiguity. The early majority buys on peer evidence and tolerates almost none. Critically, they weight sources differently: the majority discounts vendors and innovators alike, and trusts the judgment of operators who resemble them. Your enthusiastic early customers will therefore not persuade the next group by being enthusiastic, since enthusiasm is exactly what the majority expects from an innovator and discounts. They persuade by being visitable.

Re-invention predicts durability

Adopters who modify a practice to fit their conditions keep it longer than those who follow it exactly. Rogers treats modification as evidence of understanding, not dilution: build for adaptation.

APPLIED

Sequencing intensification before automation maximizes complexity, destroys trialability, and puts the payoff behind the risky part, with a leak in the middle where a grower who succeeds no longer needs you. Collapsing both into one bought outcome, entered on a single bay for a single cycle, is what his framework recommends. It also moves the pitch from evidence to risk: prove nothing, make trying cost almost nothing.

Rural Development: Putting the Last First

Robert Chambers, 1983

> Expert advice systematically fails independent operators because the expert sees a biased sample and defines the problem from outside it.

Chambers writes about agricultural extension in developing economies, but the mechanism generalizes to any situation where the person being advised carries the full downside and the advisor carries none. His central complaint is a loop: outsiders arrive with a solution, deliver it to whoever is easiest to reach, observe results under conditions unlike the typical operator's, conclude the practice works, and interpret the resulting non-adoption as backwardness rather than as evidence about their own method.

The book is polemical and the register is academic, but the empirical core has held up across four decades of extension research, and it is closer to your actual problem than anything published by a business press.

Outsider bias has a specific and predictable shape

Advisors see operations near roads, operators who share their vocabulary and want to engage, sites visited during the good season, and enterprises already prosperous enough to host a visit. Every one of those filters selects for adoption success. You then generalize from a sample of the already-convinced, build confidence, and are genuinely surprised when the wider population does not follow. The correction is not more data, it is deliberately seeking the operators who did not respond and finding out why.

The operator defines the problem, not the advisor

Chambers inverts the standard relationship. The farmer is not the recipient of a recommendation; they hold the most information about their own binding constraint and are the only party able to evaluate whether a solution fits. Advice arriving without a diagnosis built from their situation is usually solving the advisor's favourite problem. In practice this means the first meeting should establish what they think is limiting them, and you should be genuinely willing to discover it is not yield.

Risk aversion is rational, not ignorance

An operator whose entire year rides on one production cycle is not being timid when they decline a change with an attractive average return and unknown variance. They are correctly weighting the tail where the season fails. This is a different objective function from yours, not a worse understanding of the same one. Any pitch that leads with upside and stays quiet about downside identifies the speaker, accurately, as someone who does not carry the risk. Naming the downside explicitly and then removing it is more persuasive than any projection.

Demonstration beats transfer of information

The most consistent finding across decades of extension work: operators adopt what they have seen work on a farm resembling their own, run by someone whose judgment they already trust. Resembling is meant concretely. Similar scale, similar climate, similar market channel, similar labour situation. A grower dismisses a result from an operation twice their size as irrelevant to them, and they are usually right to, because the constraints that would bind them were absent there.

Participation changes ownership of the result

When operators help design or run a trial, they defend the outcome instead of questioning it, and they explain it to peers in their own terms. A demonstration site that a group of growers helped set up generates a categorically different quality of evidence from one you built and then invited them to admire. The second reads as a showroom.

Watch for the elite capture pattern

The best-resourced operators adopt first, absorb the available support, and their results become the published proof. The mainstream observes this and correctly concludes the practice requires capital, staff or facilities they lack. If your visible wins all come from well-capitalized farms, that evidence will actively slow the majority rather than reassure them. Chambers' recommendation is to deliberately over-invest in making it work on an unremarkable operation, even at worse unit economics, because that is the case that travels.

APPLIED

This is the strongest available argument for reference farms that are ordinary rather than impressive: a visitable operation at a scale your target grower recognizes, ideally one that was struggling beforehand, beats a flagship result. It also carries an uncomfortable corollary about Ferme Décembre. As proof, your own farm is the weakest instrument you have, because you are simultaneously the outsider who designed the method and the operator reporting the outcome. It is invaluable as a development site and easily discounted as evidence.

Crossing the Chasm

Geoffrey Moore, 1991 · revised 2014

> The gap between visionary buyers and pragmatic buyers is a discontinuity, not a slope, because the two groups buy for reasons that do not translate.

Moore takes Rogers' adoption curve and argues the commercially interesting part is a specific, repeatable failure. A product sells well to enthusiasts and visionaries, growth then flattens, and the company diagnoses a volume problem and responds with more reach and more messaging. It is not a volume problem. Visionaries are buying a change they want to lead and will tolerate an incomplete product to get there first. Pragmatists are buying a proven fix to a problem they already have, and their governing rule is that they will not be first.

The examples are dated and the vocabulary is enterprise software, but the underlying mechanism is Rogers' peer-observability finding turned into a sequencing strategy, and it survives translation into agriculture without much damage.

Pragmatists buy references, and only from their own segment

The pragmatist's decision procedure is to locate someone materially like them who has already done it and survived. Not similar in outlook, similar in circumstance. This creates the trap: they require a reference from inside their segment, and you have none until one of them goes first. Selling harder does not resolve a chicken-and-egg problem, and effort spent on persuasion here is largely wasted.

The beachhead solves the trap by shrinking the pond

Pick one narrow, self-aware segment and take it completely rather than accepting scattered business across five. Narrow enough that members talk to one another, share a specific and urgent problem, and can be reached without a general campaign. Once a handful of visible operators inside that group have adopted, the reference condition is satisfied for everyone else in it at once, and the segment tips rather than converting one at a time. The counterintuitive part is that a smaller target produces faster absolute growth.

Bowling pins

The beachhead is chosen partly for what it knocks over next. Win segment one, then move to an adjacent segment where either the problem or the reference carries across. Sequence matters as much as selection, and a segment that is easy to win but connects to nothing is a dead end.

The whole product

Visionaries assemble missing pieces themselves and enjoy it. Pragmatists expect their problem solved end to end, including the unglamorous parts: onboarding, training, integration with how they already work, and a person to reach when something fails mid-season. Anything absent is not a product gap to them, it is a reason to wait another year. Moore's point is that the whole product is defined by the segment, so a narrow beachhead also makes completeness affordable.

A compelling reason to buy now

Pragmatists move on pain with a deadline, not on opportunity. General upside makes them interested, and interested is a stable state they can occupy indefinitely. Something must make this season the season, and if the offer does not contain that trigger, the sale reliably slips a year at a time without ever being lost.

The chasm looks like a marketing problem and is not

The reflex when growth stalls is broader reach, more content, thinner effort per prospect. Moore argues for the opposite: concentrate past the point of comfort, accept a smaller addressable market for a period, and treat the spending as buying references rather than buying awareness. Most companies cannot bring themselves to decline the adjacent revenue, which is why the chasm claims them.

APPLIED

Your beachhead has a plausible definition already: growers who are space-constrained rather than time-constrained, with demand they cannot fill and no realistic path to expanding the structure. They have a specific urgent problem, they recognize themselves instantly in the description, and yield per square metre is the only lever available to them, which makes your argument non-optional rather than interesting. That segment also supplies the buy-now trigger you currently lack, because a season of unmet demand has a date attached to it.

Obviously Awesome

April Dunford, 2019

> Positioning is choosing the context in which your product gets evaluated. The buyer supplies a context whether or not you do, and theirs is usually wrong for you.

Dunford's practical claim is that most weak positioning is inherited rather than chosen. Companies describe themselves in terms of what they built, or in terms of the category they came from, and silently inherit that category's assumptions about what the thing does, who buys it, how it is evaluated and what it should cost. Her method runs backwards: establish what is genuinely different first, and only then select the frame that makes the difference matter.

The book is short and unusually operational. It is also the best available check on the instinct to declare a new category, which is the instinct most founders in your position have.

Start with competitive alternatives, including doing nothing

Ask what the buyer would do if you did not exist. For your market that is roughly: hire an agronomist a few times a season, read and experiment alone, or simply carry on as they are. Doing nothing is almost always the strongest competitor, and treating it as absence of demand rather than as an active alternative is the most common error. It has real advantages that deserve to be named: it is free, familiar, and carries no execution risk.

Isolate unique attributes, then translate them into value

Attributes are things you have that the alternatives do not. Value is what those enable for the buyer. The discipline is refusing to claim value your alternatives also deliver, however true the claim is. Anything a consultant can also say is not positioning, it is table stakes, and including it makes the genuinely distinct claims harder to hear.

Find who cares a lot

Identify the characteristics that make a buyer care disproportionately about your specific value. This is Moore's beachhead reached from the positioning side rather than the sales side, and the two converging on the same segment is a good sign you have found it.

Then, and only then, choose the market frame

The category you claim tells the buyer what to compare you against, which questions to ask, and what a reasonable price looks like. Choose the frame in which your unique value is obviously important, rather than the one that describes you most accurately. Accuracy is not the objective; a correct description inside a frame that makes your advantage look marginal is worse than a slightly awkward frame that makes it central.

Creating a new category is a real option and an expensive one

Her most useful contrarian point, and the one to weigh hardest. A new category means funding the buyer's education yourself, and it only pays when no existing frame can hold your value. Most companies who believe they need a new category actually need a better-chosen existing one, entered first and bent afterward. Entering an established frame and then shifting it from inside is slower to feel heroic and much cheaper to execute.

Trend is a modifier, never a foundation

You can attach to a movement to make the value feel current, but only layered on a solid frame. Trend alone attracts attention from people who will never buy, and it dates.

> APPLIED
>
> "Greenhouse consultant" is the inherited frame and it damages you twice. It sets the comparison against agronomists who sell advice, which caps your price near their day rate regardless of the value delivered. And it tells a grower protecting a forty-hour week that they will receive recommendations they must then execute, which parses as more work rather than less. You described the real product as doing the daily work through technology. Growers already have a frame for paying someone to do work they would otherwise do themselves, and that frame carries no education cost, invites no hourly comparison, and prices on outcome.

The 22 Immutable Laws of Marketing

Al Ries & Jack Trout, 1993 · read alongside Play Bigger, 2016

> Buyers hold categories in their heads as slots with an owner. Being first into a slot beats being better inside someone else's.

Ries and Trout are aphoristic, frequently overstated, and still the sharpest statement of why category matters more than quality. Their model is that a buyer's mind holds a short ranked list per category, that the list is very hard to reorder once formed, and that most marketing budget is spent fighting for second place inside a category a competitor defined. Play Bigger makes the same argument at book length with contemporary cases and a category-design vocabulary; it is more actionable in places and much more padded.

Treat the laws as strong heuristics rather than laws. Several have obvious counterexamples. The value is in the sharpness, which forces a decision where more balanced books let you defer.

The Law of the Category

If you cannot be first in a category, define a new category you can be first in. The operative word is define. Categories are not discovered in the market, they are asserted by someone and then either taken up or ignored, and the assertion costs nothing to make and a great deal to sustain.

The Law of the Mind

Being first into the buyer's head beats being first into the market. Several category owners arrived years late and simply named the thing clearly before anyone else bothered to. This is the law that makes category creation tempting, and it is true as far as it goes; what it omits is the cost of the naming campaign.

The Law of Focus

The strongest available position is ownership of a single word or idea. Ownership requires three things at once: the word must be unclaimed, it must map to a decision the buyer actually makes, and it must be repeated long past the point where you personally are tired of it. Most companies satisfy the first two and quit on the third.

The Law of Sacrifice

To own a word you must give up other words, other segments and other lines of business. Their claim, which matches what usually happens, is that positioning failures are more often failures of nerve than of insight. The company can see the narrow position clearly and cannot bring itself to decline the adjacent revenue, so the position never consolidates.

The Law of the Opposite

Going after an established alternative, do not be a better version of it. Be the opposite, built specifically on what its strength forces it to be bad at. This is Helmer's counter-positioning stated in plainer language, and it is the most immediately usable law in the book because it identifies a position the incumbent cannot follow you into without dismantling their own economics.

The Law of Line Extension

The most frequently violated law in the set. Stretching one name across more offerings dilutes the association that made it valuable in the first place, and the dilution is invisible in the short run because the early extensions sell on borrowed strength.

APPLIED

Set against Dunford, this is a genuine tension worth sitting with rather than resolving quickly. Ries and Trout say assert the new category; Dunford says borrow an existing frame and bend it. The reconciliation is situational: category creation pays when the buyer has no adequate frame and you can afford to educate them, and neither condition clearly holds for you yet. The Law of the Opposite is what to take now, because it costs nothing and is structurally defensible. Your identity is the inverse of the consultant, someone who does the work rather than advising on it, and that is precisely the move a consultant's billing model prevents them from copying.

Switch: How to Change Things When Change Is Hard

Chip & Dan Heath, 2010

> When someone agrees with you and still does not act, the missing ingredient is almost never persuasion.

The Heaths organize change around a metaphor borrowed from Jonathan Haidt: a rider on an elephant, travelling a path. The rider is deliberate reasoning, good at analysis and prone to overthinking and stalling. The elephant is emotion and habit, which supplies all the actual energy and cannot be argued with. The path is the surrounding environment. Most change efforts address one of the three and neglect the others, and what gets labelled resistance is usually one of three separately diagnosable failures.

This is the most directly operational book in the set, and the one whose recommendations you could apply to next season's onboarding without any strategic decisions being made first.

What looks like resistance is often a lack of clarity

The rider needs specific instruction, not principles, and will generate objections indefinitely when given the latter. "Increase planting density and manage the additional vigour" is a principle. "Plant at this spacing, prune to this number of stems, and I will tell you when to adjust" is an instruction. Ambiguity produces paralysis that is very easily misread as reluctance, and the two call for opposite responses.

What looks like laziness is often exhaustion

Self-control behaves like a depletable resource, and change consumes it continuously until the new behaviour becomes routine. Your buyers are already at capacity by definition, since running out of hours is the problem they are trying to solve. Any change requiring sustained willpower across a full season will lose to the existing routine no matter how convinced they were in the meeting, and the failure will arrive quietly in week six rather than as a stated objection.

Find the bright spots

Instead of analyzing why the general population fails, study the few already succeeding under the same constraints and copy exactly what they do. Bright spots are more actionable than root causes, cheaper to investigate, and far more credible to peers, because the answer arrives as "someone like you already solved this" rather than as an expert's theory.

Script the critical moves

Reduce the change to a small number of unmistakable behaviours with timing attached. The value is in eliminating decisions rather than in the content of the advice, and a slightly suboptimal protocol that is unambiguous will outperform an optimal one that requires judgment calls from someone who is already overloaded.

Shrink the change

Make the first step small enough that failing it costs nothing, and arrange for visible progress quickly enough that the elephant is given a reason to continue. This is Rogers' trialability derived independently from psychology, which is a reasonable indication both are describing something real.

Tweak the environment

Changing the situation is usually cheaper and far more reliable than changing the person. If the correct action is the default and the incorrect one requires effort, motivation stops being the binding constraint entirely. Environment change also survives the adopter's bad weeks, which motivation does not.

APPLIED

Your product sits unusually well here, because software that runs daily decisions *is* environment change rather than motivation change. That is a stronger and more specific claim than time saved, and it is one a consultant structurally cannot make, since advice always terminates in a decision the grower still has to remember to make. Worth saying plainly: not "you will save hours" but "these decisions stop being yours to remember." It also suggests the onboarding should script a small number of dated moves rather than teaching a method.

Monetizing Innovation

Madhavan Ramanujam & Georg Tacke, 2016

> Willingness to pay is a product design input, not a launch decision. Establish it before you finish building.

A thin book built around one idea, but the idea is correct and well-evidenced. Their observation from pricing consulting is that most product failures are monetization failures that happened long before anyone discussed pricing: the team built to a specification, then asked what to charge, and discovered the value created was not value anyone would pay for at the cost of delivering it. By then the specification is sunk and the options are all bad.

Skim the book, take the discipline seriously. The willingness-to-pay conversation is the cheapest research available to you and the one most consistently avoided.

Have the pricing conversation early and directly

Ask what someone would pay while the product is still changeable, so the answer can influence what gets built rather than merely being absorbed. Their point that the conversation feels uncomfortable is fair and beside the point: the discomfort is the information. A prospect who enjoys discussing your product and becomes evasive about price has told you something a survey would not.

Segment by willingness to pay, not by description

Segments defined by farm size, region or crop are convenient and usually not the real fault lines. Useful segmentation groups buyers by how much your specific value is worth to them, which frequently cuts across the obvious categories and produces groups that look heterogeneous on paper. Space-constrained versus land-rich is a willingness-to-pay segmentation, and it is the one that matters here. Small versus large is not, and organizing around it will hide the pattern.

Four ways monetization fails

The feature shock, loaded with capability nobody pays extra for. The minivation, a real improvement priced and positioned too modestly to matter. The hidden gem, genuine value the company never brings to market properly, usually because it emerged as a side effect. And the undead, a product launched into a segment that never wanted it. Each has a different remedy, and misdiagnosing which one you have costs a full cycle, so the taxonomy earns its place.

Choose a pricing model, not merely a number

Subscription, outcome-based, per unit of production, or tiered access are strategic choices that determine who buys and how they judge whether it was worth it. For an outcome product, charging in a unit the buyer already thinks in makes the value argument almost automatic. Kilograms, bays and beds are units your buyer uses daily. Hours are the consultant's unit, and pricing in hours silently invites the consultant comparison no matter how you describe yourself elsewhere.

Behavioural pricing tactics are real and strictly secondary

Anchoring, bundling and good-better-best all work and are worth using. They amplify a sound value story and cannot substitute for one, and reaching for them early is usually a sign the underlying willingness to pay was never established.

APPLIED

The unit question is the live one and it connects straight back to positioning. Pricing per hour or per consultation locks you inside the frame you are trying to escape, and no amount of messaging will pull you out of it while the invoice says otherwise. Pricing against the outcome, whether that is a share of the yield gain, a rate per square metre under management, or a flat fee per bay per season, makes the comparison your customer runs the correct one and makes the value legible without a spreadsheet. It also makes the offer trialable in Rogers' sense, since one bay is both a price point and a pilot.

SYNTHESIS

What the Seven Agree On

Where the books converge, and the one place they conflict

One purchase, not two. Rogers on complexity, Moore on the whole product, and the Heaths on shrinking the change all point the same direction. Intensification and automation sold in sequence is two adoption decisions with the payoff sitting behind the risky one, and a leak in the middle where a grower who succeeds no longer needs you. Sell a yield outcome that the technology delivers.

Narrow to the point of discomfort. Moore's beachhead, Dunford's who-cares-a-lot, the Law of Sacrifice and willingness-to-pay segmentation converge on one target. Space-constrained growers with unfilled demand and no room to expand are the population for whom yield per square metre is the only available lever, which makes your argument non-optional rather than merely interesting.

Visitable proof beats documented proof. Rogers on observability, Chambers on demonstration, Moore on references: three vocabularies for one finding, and it is the finding with the most money in it. Two or three ordinary farms at recognizable scale, walkable by a skeptical grower an hour's drive away, outweigh any volume of case studies. Chambers adds the uncomfortable corollary that your own farm is the weakest of these, because you are the outsider who designed the method.

The frame decides the price. Dunford and the pricing book arrive at this from opposite ends. Consultant framing invites comparison against a day rate and implies the grower still does the work. An outcome frame, priced in the grower's own units, removes both problems in a single move.

Where they conflict. Ries and Trout say assert a new category and own it. Dunford says a new category is expensive and you should borrow an existing frame first. The disagreement is real and the resolution is situational: category creation pays when the buyer has no adequate frame and you can afford to educate them, and neither condition looks satisfied yet. Take the Law of the Opposite in the meantime. It costs nothing and it is defensible.

The caveat that outranks all of it. None of these books address the case where the founder's attention is the binding constraint. Every recommendation above assumes sustained effort against one narrow front across several production cycles, which is exactly the resource that running several ventures simultaneously removes. Before choosing among these, it is worth testing whether the crux sits inside this business or one level above it.